

Amazon Redshift Advisor - AWS tool that provides specific recommendations to increase performance and reduce costs associated with Amazon Redshift

Archive - Data that is stored in AWS Glacier

Asynchronous decoupling - An approach to decoupled components in which the components do not need to be present at all times for the system to function

Aurora - A MySQL and PostgreSQL-compatible relational database engine that combines the speed and availability of high-end commercial databases with the simplicity and cost-effectiveness of open-source databases

Auto Scaling - A feature that helps you maintain application availability and allows you to scale your Amazon EC2 capacity automatically according to conditions that you define

AWS Batch - A service for efficiently running batch jobs against AWS resources

AWS Cloud Compliance - A set of online tools to assist you in achieving regulatory compliance with your AWS solutions

AWS Config - A service that enables you to assess, audit, and evaluate the configurations of your AWS resources

AWS Global Infrastructure - Regions and Availability Zones located around the globe

AWS KMS - The encryption management service for AWS

AWS Simple Monthly Calculator - An online tool that allows you to estimate monthly AWS charges

AWS Storage Gateway - A service that seamlessly enables hybrid storage. between on-premises storage environments and the AWS Cloud

AWS X-Ray - A tool for tracking execution of code in distributed applications

Billing metrics - A set of metrics in CloudWatch that permit the careful monitoring of AWS costs; you can also set alarms based on this information

Bucket - Logical storage container in S3

Bursting Throughput Mode - Throughput scales as the file system grows

CapEx - Capital expenditures

ClassicLink - A feature that allows you to connect EC2-Classic instances to your VPC

CloudFormation - A service that gives developers and system administrators an easy way to create and manage a collection of related AWS resources, provisioning and updating them in an orderly and predictable fashion

CloudFront - A global content delivery network (CDN) service that accelerates delivery of your websites, APIs, video content, or other web assets

CloudTrail - A web service that records AWS API calls for your account and delivers log files to you

CloudWatch - A monitoring service for AWS Cloud resources and the applications you run on AWS

Convertible RI - A reserved instance that permits modifications from the original reservation

Cooldown period - Mandatory waiting period between Auto Scaling actions

Cost Explorer - A tool in the Management Console that permits you to create detailed reports based on your AWS costs

Data tier - The storage media for the data required by an application

Database Migration Service - A service that helps you migrate databases into or out of AWS easily and securely

DAX DynamoDB Accelerator – A caching engine specially designed for DynamoDB

Decoupling - Creating components that have the capability to be autonomous from each other to some degree

DHCP option sets - Components that allow you to provide settings such as a DNS name or DNS server address to instances

Direct Connect - A tool that makes it easy to establish a dedicated network connection from your premises to AWS

Directory Services - A service that enables your directory-aware workloads and AWS resources to use managed Active Directory in the AWS Cloud

Distribution - A group of resources cached and shared by CloudFront

DNS Domain Name System - DNS server provided by AWS or the one or more servers that you provide

Durability Resiliency - In data storage, this term means the chance of data loss

DynamoDB - A fast and flexible NoSQL database service for all applications that need consistent, single-digit millisecond latency at any scale

EBS-optimized instance - Dedicated bandwidth between EC2 and EBS

Edge location - A special data center designed for delivering CloudFront cached resources

Egress-only Internet gateway - A VPC component that provides Internet access for IPv6 addressed instances

Elastic Beanstalk - An easy-to-use service for deploying and scaling web applications and services

Elastic Block Store - A resource that provides persistent block storage volumes for use with EC2 instances in the AWS Cloud

Elastic Compute Cloud (EC2) - Virtual machine compute resources in AWS

Elastic Container Registry - A fully-managed Docker container registry in AWS

Elastic Container Service - A scalable, high-performance container management service that supports Docker containers

Elastic Container Service for Kubernetes - A fully managed service that makes it easy for you to use Kubernetes on AWS without having to be an expert in managing Kubernetes clusters

Elastic File System - A service that provides simple, scalable file storage for use with Amazon EC2 instances in the AWS Cloud

Elastic IP addresses - Public IP addresses that you can assign flexibly in your AWS VPC

Elastic Load Balancing - A feature that automatically distributes incoming application traffic across multiple EC2 instances

ElastiCache - A web service that makes it easy to deploy, operate, and scale an in-memory cache in the cloud

Elasticity - The ability of resources to scale up and down

Fargate - A technology for Amazon ECS and EKS that allows you to run containers without having to manage servers or clusters

FT Fault tolerance – A subset method of high availability designed to create a zero downtime solution and zero performance degradation due to a component failure

Glacier - A secure, durable, and extremely low-cost storage service for data archiving and long-term backup

Global secondary index - An index with a partition key and a sort key that can be different from those on the base table

Greengrass - An AWS service that allows IoT devices to run code that is synched from the cloud

Groups - Objects that contain user accounts; permissions can be assigned to groups and not users to implement scalability in the design

HA High availability - A goal in solution design that focuses on automatic failover and recoverability from an issue in the solution

HTTPS - The secure version of HTTP for protecting data in transit

Identity and Access Management (IAM) - A tool that enables you to securely control access to AWS services and resources for your users

Initialization - A process in which each block of an EBS volume is accessed in advance of deployment; it is used when you have restored a volume from a Snapshot

Instance store - Legacy ephemeral storage option for EC2 instances

Internet gateway - Virtual device that provides Internet access to instances

Key Management Service (KMS) - A managed service that makes it easy for you to create and control the encryption keys used to encrypt your data

Kinesis - A data platform that makes it easy to collect, process, and analyze realtime, streaming data

Lambda - A serverless compute service of AWS that permits the execution of custom code or functions without the provisioning of virtual servers

Latency - Delay (from the storage to the instance, for example)

Lazy loading - Caching data strictly based on requests for the data

Lightsail - An interface that enables you to launch and manage a virtual private server with AWS

Local secondary index - An index that has the same partition key as the base table but a different sort key

Logic tier - The server that contains the code providing the functionality of the application or technology

Loose decoupling - Creating decoupled components that have few dependencies on each other

Memcached - A simple caching engine supported by ElastiCache

Multi-AZ - The reference to using multiple Availability Zones in an AWS design; often done to ensure high availability (HA)

Multi-Tier architecture - The use of multiple systems or platforms to implement an application or technology

NAT - Network Address Translation as supported in AWS

Network ACL - A security mechanism in AWS used to control traffic flow in and out of subnets within VPCs

Network interface - Virtual interface that represents a NIC

NFS Network File System - a standard shared file system protocol

NIST - The National Institute of Standards and Technology, which provides excellent guidance on best practices

Object storage - The data storage technique of AWS S3

OpEx - Operating expenditures

OpsWorks - A configuration management service that uses Chef, an automation platform that treats server configurations as code

Patch Manager - A component of AWS Systems Manager that automates the deployment of patches to operating systems in your AWS infrastructure

Policy - An object in IAM that defines the actual permissions assigned to a resource or service in AWS

Presentation tier - Components that users interact with for a user interface

Provisioned Throughput Mode - An operating mode available for high throughput needs or requirements greater than Bursting Throughput mode allows

Redis - A cache engine supported by ElastiCache that supports complex data types and multi-AZ deployments

Redshift - A fast, fully managed, petabyte-scale data warehouse that makes it simple and cost-effective to analyze all your data using your existing business intelligence tools

Relational Database Service (RDS) - A service that makes it easy to set up, operate, and scale a relational database in the cloud

Reserved instances - AWS capacity utilized at a discount over default on-demand pricing as a result of the reservation

Role - Similar to a user but does not represent one user in the organization; instead, it is intended to be assumable by anyone who needs it

Root user - Role created when you create your AWS account; this role has unrestricted access to all aspects of your account

Route 53 - A highly available and scalable cloud Domain Name System (DNS) web service

Route table - Routing instructions for your subnets

RPO - Recovery point objective; the point to which data must be restored

RTO - Recovery time objective; the amount of time in which a recovery must be completed

Scheduled RI - Reserved instances that will operate within a certain time window

Security group - A security component in AWS used to secure EC2 and other resources

Serverless Application Repository - A service that enables you to deploy code samples, components, and complete applications for common use cases such as web and mobile back ends, event and data processing, logging, monitoring, IoT, and more

Simple Notification Service (SNS) - A fast, flexible, distributed, and fully managed push notification service that lets you send individual messages or to distribute messages to many recipients

Simple Queue Service (SQS) - A fast, reliable, scalable, distributed, and fully managed message queuing service

Simple Storage Service (S3) - Object storage with a simple web service interface to store and retrieve any amount of data from anywhere on the web

Single-Tier Architecture - A type of architecture in which all required components for a software application or technology are provided on a single server or platform

Snowball - A petabyte-scale data transport solution that uses secure appliances to transfer large amounts of data into and out of AWS

Standard RI - A standard reserved instance; this type provides the greatest potential cost savings

Storage classes - Different classes of S3 storage used to balance performance and costs

Synchronous decoupling - An approach to decoupling that means decoupled components must be present for the system to function properly

TCP Selective Acknowledgment - A performant S3 service that improves recovery time after large numbers of packet loss

TCP Window Scaling - A performant S3 service that supports window sizes larger than 64 KB

Technical Account Managers (TAMs) - Experts that are available to you if you have an Enterprise Support level of service

Trusted Advisor - An online resource to help you reduce cost, increase performance, and improve security by optimizing your AWS environment

TTL - The time to live value in Route 53 that dictates the amount of time resource record information cached by DNS resolvers

Users - Components in IAM that represent users in your organization

Vault - Logical storage structure in AWS Glacier

Versioning - Keeping multiple copies of objects in AWS S3 for various time points

Virtual Private Cloud (VPC) - Virtual network components in AWS

Volume queue length - The number of pending I/O requests for a device

VPC endpoints - Devices that allow you to privately connect your VPC to supported endpoint services

VPC peering - A connection that permits communications between VPCs

Web Application Firewall (WAF) - A feature that helps protect your web applications from common web exploits that could affect application availability, compromise security, or consume excessive resources

Write through - A caching strategy designed to minimize stale cache data